

MODUL UNIVERSITY VIENNA


SEID 2364

SESSION 11

Applied Ethical Analysis

David B. Smith

*Misinformation, Bad
Actors, and Professional
Responsibility*



Establish continuity from the prior weeks. This session does not introduce ethics from scratch. Students have already worked with ethical theory, the EU AI Act, mediation pathways, risk, distance, authority, trust, and EDOCA. Today extends that work by asking how ethical responsibility changes when systems include uncertainty, misinformation, bad actors, professional obligations, and conflicting frameworks.

Today's Shift

From evaluating systems



to evaluating how systems become harmful

From ethical theory



to professional responsibility

From unintended harm



to bad actors and preventable negligence

Frame the conceptual shift. Earlier weeks focused heavily on bias, mediation, distance, observability, capability, and trust. This week asks what happens when some actors intentionally exploit systems, or when professionals become harmful through negligence, recklessness, or institutional incentives.

Case Anchor

Shared Case: COVID-19 Misinformation Moderation

- 1 Read the case card.
- 2 Work individually in your private journal — do not discuss yet.
- 3 Do not decide yet whether moderation was ethical or unethical.
- 4 Identify what your judgment depends on.
- 5 Focus on your own reasoning.

Use only the short COVID case card here. Do not introduce ACM, IEEE, EU governance, or the expanded case study yet. Students should not optimize their responses for correctness. They should record their initial assumptions. Remind them that private journal responses are not being graded directly and are used to surface current thinking.

PROVOCATION 1

**Do you think this misinformation
moderation was ethically justified?
Why or why not?**

Write a short informal essay.

Identify what your answer depends on.

Students answer individually before discussion. Strong responses should identify conditions rather than jump to yes/no conclusions. Look for criteria such as public safety, free expression, scientific uncertainty, authority, transparency, appeal, proportionality, and trust.

PROVOCATION 2

Imagine three versions of the same moderation pathway:

1

The platform acts under its own policy.

2

A public-health authority flags content and the platform acts.

3

A regulator, court, or government body requires action.

How did your ethical evaluation change across these scenarios — and what drove that change?

Separate action from authority. The same visible outcome—labeling, downranking, removal, or suspension—may have a different ethical interpretation depending on who initiated it, who authorized it, and who controls the pathway.

Ask students to name the variable that changed. Did the system change, or did the context change? Did the moral weight of harm change? Did public authority become more legitimate or more dangerous? Did responsibility become clearer or more distributed?

PROVOCATION 3

Identify the different agents in this system.

For each agent, classify their access using the three dimensions we have been discussing:

Observability • Control • Authority

Where are these aligned — and where are they separated across different agents?

Students identify the agents themselves — do not seed them. The slide references the three access dimensions (Observability, Control, Authority) that students have been working with. Key nuance to draw out in discussion: individual posters have limited control — they can place content into the space, but how it propagates is governed by other agents (algorithmic systems, platform moderation, institutional authority). The algorithm has enormous control over signal propagation but ambiguous authority. The platform has operational control but may be acting under external authority. Public-health authorities may have expertise and legitimacy but limited direct control. Students discussed these complexities last week, so they have the vocabulary. The core question is where Observability, Control, and Authority align and where they are separated across different agents.

Observability — can you see what the agent is doing and why?

Control — who directs or constrains the agent's actions?

Authority — what gives the agent the right to act?

PROVOCATION 4

For each of the three moderation pathways from Provocation 2, what types of harm emerge — and how severe are they?

Consider harm to individuals, communities, institutions, and society.
Rank the types and severity of harm you identify.

This provocation explicitly ties back to the three moderation pathways from Provocation 2: (1) platform acts under its own policy, (2) public-health authority flags and platform acts, (3) regulator/government requires action. Students now analyze harm across all three. The slide prompts them to consider harm at multiple scales: individuals, communities, institutions, and society. Society must appear as a stakeholder — if students omit it, prompt them. Societal harm includes erosion of public trust, democratic legitimacy, epistemic commons, and social cohesion. Strong responses will distinguish between types of harm (e.g., speech suppression, trust erosion, health risk, institutional overreach, democratic damage) and recognize that the same moderation action produces different harm profiles depending on which pathway governs it. The Observability/Control/Authority mapping from Provocation 3 should inform this analysis directly.

DISCUSSION ACTIVITY

From Agents to Harms: A Collaborative Analysis

1. Identify the Agents

Who acts in this system? Look beyond the obvious three. Consider algorithms, audiences, economic structures, and institutional incentives.

2. Explore the Harms

For each agent: What harms can emerge from their actions? Are they acting from negligence, recklessness, or strategic intent? What does mitigation look like?

3. Capture and Build

Document your analysis in a shared resource. Map each agent against Observability, Control, and Authority. This becomes a living reference we will return to.

The most dangerous agents are often the ones nobody thinks to name.

Professional Frameworks and Codes of Ethics

How do formal frameworks structure the ethical analysis you just performed?

Section pivot. Students have just: (1) performed ethical analysis using their own vocabulary, (2) mapped agents, authority, and harm across a real system, and (3) found places where their analysis needed more structure. Consider briefly asking two or three students what surprised them or where they felt uncertain. Use their responses to motivate why professional codes and governance frameworks exist: they formalize the kind of reasoning students just did, and they make disagreements between frameworks visible. The lecture now shifts to ACM, IEEE, EU governance, and BBS/EDOCA as formal lenses for the same kind of work.



PURPOSE: This slide pivots from COVID as a single case study to the broader landscape. The point is that the analytical structures students just applied (agents, authority, harm, observability, control) transfer directly to every domain shown here. Walk through each domain on the slide, pausing to discuss.

ELECTIONS: Voter suppression narratives, fabricated candidate statements, deepfake endorsements. Authority is fragmented — election commissions, platforms, media, and political parties all claim different roles. Ask: who has observability over what voters actually see?

CLIMATE COMMUNICATION: Decades of coordinated doubt manufacturing. Fossil fuel industry funded counter-narratives that mimicked scientific discourse. Key tension: industry actors had control over messaging channels while scientific institutions had authority but limited reach. Distortion operated over very long timeframes.

WAR REPORTING: Fog of war meets information warfare. State actors, militias, and citizen journalists all produce conflicting accounts. Observability is extremely limited — verification is slow while narratives spread instantly. Discuss how authority to narrate shifts between embedded journalists, government sources, and social media witnesses.

FINANCIAL PANIC: Rumours move markets before regulators can respond. Social media-driven trading (e.g. meme stocks) demonstrated how control over financial narratives shifted from institutional analysts to decentralised online

communities. Ask: does the speed of information flow itself become a form of distortion?

CIVIL UNREST: Protest movements generate competing narratives in real time. Selective footage, false flag accusations, agent provocateur claims. Both state actors and activist groups attempt to control the narrative. Observability is contested — body cameras, bystander video, and official statements often contradict each other. Who has authority to define what happened?

DISASTER RESPONSE: Misinformation surges during crises when people are most vulnerable. False evacuation routes, fake charity scams, conspiracy theories about causes. Emergency management agencies have formal authority but limited control over social media channels where misinformation spreads fastest.

SCIENTIFIC CONTROVERSIES: Vaccine hesitancy, GMO debates, 5G conspiracy theories. Scientific consensus exists but competes with emotionally compelling counter-narratives. The energy/effort asymmetry is stark: producing rigorous science takes years; producing doubt takes a blog post. Peer review confers authority, but social media algorithms do not weight authority.

AI-GENERATED MEDIA: The newest domain. Deepfakes, synthetic text, AI-generated images collapse the energy/effort dimension — what once required a studio now requires a prompt. Observability is under direct threat because detection lags behind generation capability. Ask: if no one can reliably distinguish real from synthetic, what happens to authority?

DISCUSSION APPROACH: Walk through domains one at a time. For each, ask students: (1) Who are the agents? (2) Where does authority sit? (3) What can and cannot be observed? (4) Who has control over the information flow? Draw explicit parallels back to the COVID case they just analysed. The goal is for students to recognise that COVID was not exceptional — it was a particularly visible instance of a recurring structural problem.

Infodemic

Too much information, including false or misleading information, during a disease outbreak.

The problem is not only false information.

The problem is overload, uncertainty, speed, trust, and authority.

Introduce WHO's term at a high level. The definition should be enough for class use; detailed citations can live in the case study. Emphasize that an infodemic includes true, false, outdated, uncertain, confusing, and deliberately manipulative information.

Mediated Information Crisis

Information moves through technical, institutional, and social systems in ways that make it difficult to determine:

What is true?

Who has authority?

What action is appropriate?

Who is responsible?

This is the generalized concept. It connects the case directly to BBS. Signals move through multiple sources, vectors, destinations, institutions, platforms, and communities. Distortion and trust failures occur across mediation pathways.

From Information Crisis to Agency

Information does not move by itself.

Actors create, share, suppress, amplify, classify, and interpret signals.

-  Some act in good faith.
-  Some act carelessly.
-  Some act strategically.
-  Some act maliciously.

Pivot from environment to agency. Misinformation is not one ethical category. The same false claim can be carried by different kinds of actors with different degrees of awareness, responsibility, and intent.

Information Disorder Operates Across Multiple Layers

Informational Distortion

Problems related to informational content, framing, context, uncertainty, and interpretation.

Systemic Distortion

Problems related to amplification systems, mediation pathways, algorithms, infrastructures, and incentives.

Harmful Agency

Problems related to how people, institutions, and computational systems behave within mediated systems.

PURPOSE: This is the master taxonomy slide. It introduces the three-layer model that replaces the simpler 'types of false information' framing. Each layer gets its own detailed slide next.

KEY SHIFT: We are no longer only talking about types of false information. We are now distinguishing among: (1) informational distortion — problems with the content itself, (2) systemic distortion — problems with how systems move, amplify, and transform information, (3) harmful agency — problems with how actors behave within mediated systems.

DELIVERY: Present this as an overview. Say: 'The following slides will unpack each of these layers.' Do not spend too long here — the detail comes next. The purpose is to establish the structure so students have a mental map.

CONNECT TO EDOCA: These three layers map onto the EDOCA chain. Informational distortion relates to the Energy/Effort and Distortion dimensions. Systemic distortion relates to the mediation pathway itself. Harmful agency relates to Observability, Control, and Authority — who can see, who can act, who has permission.

False Information

The baseline category — information that is simply wrong.

DEFINITION

Factually incorrect information. The broad descriptive umbrella — no assumption about intent, motive, or origin.

EXAMPLES

A news article misreports a statistic due to a data error. A student submits an essay with an incorrect historical date. A Wikipedia entry contains an outdated figure.

KEY QUESTION

If no one intended to deceive, does it still matter? When does inaccuracy become an ethical problem rather than just an error?



PURPOSE: This slide introduces the broadest, least loaded term in the taxonomy. False information is purely descriptive — it says nothing about who created it, why, or how it spread. The point is to establish that not all false claims carry the same ethical weight, and that we need more precise language to distinguish between them.

KEY POINTS TO MAKE:

- False information is a factual description, not a moral judgement. It simply means the claim does not correspond to reality.
- This is the umbrella under which misinformation, disinformation, and malinformation all sit — but each adds a layer of intent or context.
- Many everyday instances of false information are mundane: typos, outdated data, honest mistakes. These are not ethically neutral (they can still cause harm), but they are ethically different from deliberate deception.
- In EDOCA terms: the distortion exists, but the agent may have minimal awareness of it. There is no intent to distort — the energy/effort behind the falsehood may be zero (it was inherited, not created).

DISCUSSION PROMPT: Ask students — if a doctor gives a patient incorrect dosage information because of a textbook error, is that doctor ethically responsible? They did not create the error, but they are the agent who transmitted it. This sets up the distinction between creating and transmitting false information, which becomes critical in later slides.

CONNECT TO EDOCA: Observability is the key dimension here. The person

sharing false information may have no way to observe that it is false. Control is limited — they cannot correct what they do not know is wrong. Authority may be high (a doctor, a teacher) which amplifies the downstream impact regardless of intent.

Misinformation

False information spread without deceptive intent.

DEFINITION

False or misleading information shared without intent to harm. The sharer believes it or has not verified it.

EXAMPLES

A grandmother shares a Facebook post claiming lemon water cures cancer because she genuinely wants to help. A colleague forwards an unverified rumour about layoffs because they are worried.

KEY QUESTION

Is good faith a defense? If you share something harmful but believed it was true, are you ethically responsible for the consequences?



PURPOSE: This slide introduces intent as a distinguishing factor. Misinformation is false, but the person sharing it does not know that — or at least does not intend to deceive. The ethical question shifts from "is this true?" to "should you have checked?"

KEY POINTS TO MAKE:

- Misinformation is the most common form of information disorder. Most people who share false claims are not bad actors — they are acting in good faith based on incomplete or incorrect information.
- The ethical tension is between intent and impact. A grandmother sharing a miracle cure post has no malicious intent, but if someone follows that advice instead of seeking medical treatment, real harm results.
- In professional contexts, the standard is higher. A journalist who publishes an unverified claim may not intend to mislead, but their professional code requires verification. Misinformation from a professional is closer to negligence.
- Social media architectures make misinformation worse: sharing is frictionless, verification is effortful, and emotional content spreads faster than factual corrections.

DISCUSSION PROMPT: Ask students — should platforms treat misinformation differently from disinformation? If the content is identical (same false claim), but one was shared in good faith and one was planted by a troll farm, should the platform's response differ? This previews the "Same Claim, Different Agency" slide coming later.

CONNECT TO EDOCA: The agent has low observability (they cannot see that the claim is false) and low control (they are an end-node, not a gatekeeper). But they may have authority in their social circle — a trusted friend, a parent, a community leader — which gives their sharing disproportionate weight. The distortion is amplified not by intent but by trust.

Disinformation

DEFINITION

False or misleading information deliberately created or used to deceive, manipulate, or cause harm.

EXAMPLES

A state-sponsored campaign creates fake social media accounts to spread fabricated stories about a political candidate. A company publishes misleading safety data to avoid a product recall.

KEY QUESTION

Does the presence of intent change the ethical calculus entirely? Is a deliberate lie always worse than an accidental one, even if the accidental one causes more harm?



PURPOSE: This is where intent enters the taxonomy explicitly. Disinformation is not just false — it is strategically false. Someone created or deployed it knowing it was misleading, with the goal of achieving a specific outcome. This is the category that most people think of when they hear "fake news," but it is actually the narrowest and most specific form of information disorder.

KEY POINTS TO MAKE:

- Disinformation requires intent. This is what distinguishes it from misinformation. The same claim can be misinformation in one person's hands and disinformation in another's — it depends on what the sharer knows and intends.
- Disinformation is often systematic: it involves planning, resources, and coordination. State-sponsored influence operations, corporate cover-ups, political propaganda campaigns. These are not individual actors making mistakes — they are organised efforts to shape belief.
- The ethical weight is heavier because the agent has full awareness. In EDOCA terms, they have high observability (they know the claim is false), high control (they choose to deploy it), and they are deliberately undermining others' observability.
- However, disinformation does not always cause the most harm. A well-funded disinformation campaign that reaches 1,000 people may cause less damage than a piece of misinformation that goes viral to 10 million. Impact and intent are separate dimensions.

DISCUSSION PROMPT: Present the scenario — a pharmaceutical company

suppresses a study showing side effects. They do not publish false data; they simply do not publish the inconvenient data. Is this disinformation? It is deliberate, it is misleading, but technically nothing they published is false. This bridges to malinformation on the next slide.

CONNECT TO EDOCA: All five dimensions are active. The agent invests significant energy/effort into creating or curating the false narrative. The distortion is intentional. They deliberately reduce others' observability while maintaining their own. They exercise control over information channels. And they often claim or exploit authority (institutional credibility, scientific credentials, journalistic platforms) to make the disinformation more effective.

Malinformation True information deployed to cause harm.

DEFINITION

True or partially true information weaponized by being used out of context, selectively framed, or deployed to cause harm.

EXAMPLES

Leaked private medical records used to discredit a public figure. A real but decades-old photograph presented as evidence of a current event. Cherry-picked crime statistics used to stigmatize a community.

KEY QUESTION

If the information is true, can sharing it still be unethical? Where is the line between transparency, privacy, and weaponization?



PURPOSE: This is the most conceptually challenging category because it breaks the assumption that "true = good." Malinformation is factually accurate — or at least based on real information — but it is used in a way that is intended to cause harm, violate privacy, or mislead through selective framing. This forces students to confront the idea that truth alone is not sufficient for ethical information practice.

KEY POINTS TO MAKE:

- Malinformation challenges the simple "just tell the truth" response to information disorder. The information is true, but its deployment is harmful. This means truth and ethics are not the same thing.
- Common forms include: doxxing (publishing someone's real address), revenge sharing of private content, selective leaking of genuine documents to serve an agenda, and cherry-picking real data to create a misleading narrative.
- Context is everything. A photograph that is real but from a different time or place is malinformation when presented as current evidence. The pixels are authentic; the framing is deceptive.
- Malinformation is particularly difficult for platforms to moderate because the content itself is not false. Fact-checkers cannot flag it as "false" — the harm is in the framing, the timing, or the intent behind sharing.

DISCUSSION PROMPT: Ask students — a journalist obtains leaked emails from a political campaign that reveal genuine corruption. Publishing them is in the public interest. But the leak was orchestrated by a foreign intelligence service to influence the election. Is publishing the emails ethical? The information is true,

the source is adversarial, and the timing is strategic. This is malinformation at its most complex.

CONNECT TO EDOCA: The distortion dimension is unusual here — the information itself is not distorted, but the context around it is. Observability is selectively controlled: the audience can see the information but not the framing decisions behind its release. The agent deploying malinformation often has high control (they chose what to release, when, and to whom) and may exploit authority (journalistic credibility, institutional access) to amplify the impact. Energy/effort may be low (the information already exists) but the strategic deployment is deliberate.

Additional Informational Conditions

These are not automatically misinformation or disinformation. Their ethical meaning depends on context, framing, and deployment.

Incomplete Information

Partial or evolving information. Common in crises and breaking news. Not false, but potentially misleading if treated as complete.

Speculative Information

Claims presented before sufficient evidence exists. Prevalent in early-stage science and emerging crises. May later prove true or false.

Decontextualized Information

Information separated from the context necessary to interpret it correctly. A true statistic stripped of its caveats. A quote removed from its surrounding argument.

Synthetic Information

Computationally generated information or media. Deepfakes, AI-generated text, synthetic images. May be accurate or deceptive depending on use.

PURPOSE: This slide extends beyond the classic four-part model. These additional conditions are critical because they represent the grey areas where moderation is hardest.

INCOMPLETE INFORMATION: Early COVID reporting was full of this. The WHO initially said there was 'no clear evidence of human-to-human transmission.' That was accurate at the time — the evidence didn't yet exist. But it was incomplete and later proved misleading. Should it have been flagged?

SPECULATIVE INFORMATION: The lab-leak hypothesis is a prime example. Early on it was speculative. Some platforms suppressed it as misinformation. Later, major health agencies acknowledged it as a legitimate hypothesis. The ethical question: how should platforms handle claims that are speculative but not yet falsifiable?

DECONTEXTUALIZED INFORMATION: A true statistic can mislead when stripped of context. '99.7% survival rate' for COVID was factually derived but omitted age stratification, long-term effects, and healthcare system capacity. The information was true; the framing was deceptive.

SYNTHETIC INFORMATION: AI-generated content collapses the Energy/Effort dimension of EDOCA. What once required a production studio now requires a prompt. Not all synthetic content is harmful — but it fundamentally changes observability: can you tell what's real?

KEY POINT: None of these are automatically misinformation or disinformation. Their ethical meaning depends on context, framing, deployment, amplification, and interpretation. This is why a simple true/false binary is insufficient.

Systemic Distortion

Distortion can emerge from the systems through which information moves — not from the information itself.

Platform Mechanics

Algorithms, ranking systems, and economic incentives optimize for engagement over accuracy — systematically amplifying sensational content and rewarding virality over quality.

Coordinated Exploitation

External actors game distribution systems through bots, trend hijacking, and coordinated inauthentic behavior — weaponizing platform mechanics at scale.

AI-Targeted Attacks

Prompt injection and data poisoning target the computational infrastructure itself — manipulating AI systems that increasingly mediate what people see.

Systemic distortion is not about what is said — it is about how the system shapes what is seen.

PURPOSE: This slide shifts the conversation from 'what kind of information is problematic' to 'how do systems make information problematic.' This is Layer 2 of the three-layer model — arguably the most important conceptual addition.

KEY INSIGHT: The problem is not always the information itself. Systems amplify, compress, reward, distort, and redistribute information. Recommendation engines, algorithms, platforms, and economic structures become ethically significant actors in their own right.

PLATFORM MECHANICS: Algorithms optimize for engagement, not truth. Content that provokes emotional reactions gets surfaced, creating systemic bias toward sensationalism and outrage — regardless of whether the content is accurate. Ranking systems and economic incentives reinforce this: virality is rewarded, and platform revenue models are structurally misaligned with information quality. The distortion is not a bug — it is an emergent property of how these systems are designed.

COORDINATED EXPLOITATION: State actors and organized campaigns do not just create false content — they game the systems that distribute it. Bot networks amplify chosen narratives, trend hijacking manipulates visibility algorithms, and coordinated inauthentic behavior exploits platform trust signals. The key point: these actors weaponize the platform mechanics described above, turning systemic vulnerabilities into attack surfaces.

AI-TARGETED ATTACKS: Prompt injection and data poisoning represent a newer

class of systemic manipulation. Rather than deceiving humans directly, these techniques target the computational infrastructure that increasingly mediates what people see. As AI systems play a larger role in content curation, summarization, and search, attacking those systems becomes a way to distort information at scale.

KEY DISTINCTION: Stress the bottom-line insight: 'Disinformation targets beliefs. Adversarial manipulation targets systems.' This is a crucial conceptual upgrade — fixing information quality alone is insufficient if the systems that distribute information are themselves compromised.

CONNECT TO EDOCA: Systemic distortion maps to the mediation pathway in EDOCA. The Distortion dimension captures how signals change as they move through mediated systems. Control and Authority become critical: who controls the amplification system? Who has authority to modify it?

Bad Actors Are Not Always Villains

Harmful behaviour may emerge through many pathways — not all of them intentional.

HARMFUL BEHAVIOURS

- Negligence
- Recklessness
- Strategic exploitation
- Institutional incentives
- Computational optimization
- Malicious intent

ACTOR TYPES

- Misinformed actors
- Negligent actors
- Strategic actors
- Institutional actors
- Computational actors
- Hybrid actors

Many harmful systems emerge without a single obvious villain. Responsibility may be distributed across pathways and infrastructures. The next slides explore how the same claim takes on different meaning depending on the agent behind it.


PURPOSE: This slide introduces Layer 3: Harmful Agency. It 'pulls back the curtain' on the stereotype of the bad actor. The key message: harmful behaviour emerges through many pathways, not all intentional.

KEY INSIGHT: Many harmful systems emerge without a single obvious villain. Computational systems may amplify harm without possessing human intent. Responsibility may be distributed across pathways and infrastructures.

BRIDGE TO NEXT SECTION: This slide is a compact summary that sets up the detailed 'Same Claim, Different Agency' sequence on the following slides. The actor types listed here (misinformed, negligent, strategic, institutional, computational, hybrid) map onto the seven detailed slides that follow.

HYBRID ACTORS: This is worth highlighting — many real-world cases involve combinations. A political campaign (strategic actor) using bot networks (computational actors) to amplify genuine grassroots frustration (misinformed actors). The agency is distributed and layered.

CONNECT TO EDOCA: Harmful agency maps to the Observability, Control, and Authority dimensions. Who can see what's happening? Who has the ability to act? Who has the permission and responsibility to act? These questions become especially complex when agency is distributed.



PURPOSE: This visual diagram ties together the three-layer model into a single mediated pathway. It shows how information flows from creation through systems to impact.

SIGNAL / INFORMATION: The starting point. A claim, a report, an image, a post. It may already carry informational distortion (Layer 1) — it might be false, incomplete, speculative, decontextualized, or synthetic.

MEDIATION SYSTEM: The infrastructure through which information moves. Platforms, algorithms, recommendation engines, search rankings, social networks. This is where systemic distortion (Layer 2) operates.

AMPLIFICATION / DISTORTION / TRANSFORMATION: What the mediation system does to the signal. It may amplify it (algorithmic boosting), distort it (compression, decontextualization), or transform it (remixing, quote-tweeting out of context). The signal that arrives at the end may be very different from what was created.

INTERPRETATION / ACTION / HARM: The endpoint. How the signal is received, interpreted, and acted upon. This is where harm materialises — in changed beliefs, changed behaviours, eroded trust, physical danger.

CONNECT TO EDOCA: This flow maps directly onto the EDOCA mediation chain. Energy/Effort at the top (creating the signal). Distortion through the middle (how the signal changes). Observability, Control, and Authority distributed across all

layers (who can see, who can act, who has permission).

DELIVERY: Use this as a synthesis slide. Walk through the flow once, then say: 'Everything we have discussed — the types of information, the systemic forces, the types of actors — sits somewhere along this pathway. The course is about understanding how to intervene responsibly at each stage.'

Same Claim, Different Agency

The same false claim may be:

- Sincerely believed
- Socially reinforced
- Negligently amplified
- Strategically exploited
- Algorithmically amplified
- Institutionally tolerated
- Deliberately manufactured

PURPOSE: This is a central slide. The key insight is that the same false claim — word for word — can carry entirely different ethical weight depending on the agency behind it. Walk through each type and discuss. The ordering is deliberate: it moves from least intentional to most intentional, which maps to increasing moral culpability.

SINCERELY BELIEVED: The person genuinely thinks the claim is true. They are not trying to deceive anyone. Example: a grandmother shares a post about a miracle cure because she truly believes it works. In EDOCA terms, this agent has minimal control over the information ecosystem and no authority — they are an end-node. Ethically, are they responsible for the downstream harm? Is good faith a defence?

SOCIALLY REINFORCED: The claim circulates within a community and gains credibility through repetition and social proof. No single person may be “the source” — the group collectively reinforces it. Example: a WhatsApp family group where everyone shares and validates the same unverified health advice. The distortion here is social: the perception of consensus substitutes for evidence. Ask students: who is the agent? Can you hold a community responsible?

NEGLIGENTLY AMPLIFIED: Someone with reach (journalist, influencer, public figure) shares the claim without verifying it. They may not believe it deeply, but they did not exercise due diligence. Example: a news outlet running an unverified claim because it drives clicks. The agent here has control (reach, audience) but failed in their professional responsibility to verify. This is where professional

codes of ethics become directly relevant — negligence implies a duty of care that was breached.

STRATEGICALLY EXPLOITED: The person or entity knows the claim is questionable but uses it because it serves their agenda. They may not have created it, but they deliberately choose to amplify it. Example: a political campaign sharing misleading statistics they know are cherry-picked. This is where intent enters the picture — the agent has both control and awareness. Ask: does exploiting existing misinformation carry the same moral weight as creating it?

ALGORITHMICALLY AMPLIFIED: No human directly decided to spread the claim — a recommendation algorithm surfaced it because it drives engagement. The algorithm has enormous control over what people see, but has no understanding of truth or harm. Example: a social media feed that promotes sensational health claims because they generate clicks. Key question: who bears responsibility — the algorithm, the engineers who built it, or the company that deployed it? In EDOCA terms, observability is very low (users cannot see why content was recommended) and authority is ambiguous.

INSTITUTIONALLY TOLERATED: An institution (platform, government, regulatory body) is aware the claim is circulating and causing harm, but chooses not to act — whether for political, economic, or strategic reasons. Example: a platform that knows anti-vaccine content violates its policies but delays enforcement to avoid political backlash. The institution has both authority and control but exercises neither. This is an omission rather than an action, but the ethical weight may be just as significant. Ask: is inaction by those with authority morally equivalent to action by those without it?

DELIBERATELY MANUFACTURED: The claim was created from scratch with the intent to deceive. This is disinformation in its purest form. Example: a state-sponsored troll farm producing fake news articles to destabilise an election. The agent has full awareness, full intent, and often significant resources. In EDOCA terms, they invest high energy/effort into creating distortion, while deliberately undermining others' observability. This is where most people's moral intuition lands hardest — but ask: does it always cause the most harm? (Sometimes algorithmic amplification of a sincere belief causes more damage than a manufactured claim that no one sees.)

DISCUSSION APPROACH: Walk through each type one at a time. For each, ask: (1) What is the agent's intent? (2) How much control do they have? (3) What is their level of authority? (4) Does the ethical responsibility match the harm caused? The key takeaway is that platforms often see only the claim itself — they cannot see the agency behind it. This creates a fundamental challenge for moderation: the same content may require different responses depending on who posted it and why, but that context is often invisible.

The Platform's Problem

A platform may observe the claim.

e.g. automated detection flags a post claiming "5G causes COVID"

It may not observe the agency behind the claim.

e.g. the same post could come from a concerned parent or a coordinated bot network

It may control visibility.

e.g. demoting a post in the feed, adding a warning label, or removing it entirely

It may not know intent.

e.g. satire, whistleblowing, and state propaganda can look identical at surface level

It may act at scale before context is clear.

e.g. automated takedowns during breaking news removing legitimate journalism alongside false claims

Connect observability and control. A system may have high operational control but low contextual observability. The system may treat sincere misinformation and strategic disinformation the same way, or fail to respond to coordinated manipulation.


Bad Actor

An agent whose behavior degrades the integrity, safety, trust, fairness, or accountability of a system.

A bad actor is not always a bad person.

Define the term operationally rather than morally. The goal is not accusation. The goal is to identify harmful agency in a system.

Bad Actor Spectrum



Briefly define each category orally.

Malicious actors intend harm or manipulation.

Strategic actors exploit systems for advantage.

Reckless actors ignore foreseeable risk.

Negligent actors fail professional standards of care.

Misinformed actors sincerely believe false claims.

Institutional bad actors create harmful incentive structures.

Procedural bad actors exploit rules while undermining their purpose.

Why This Matters

Systems must be designed for:

**Good-faith
ambiguity**

and


**Bad-faith
exploitation**

Bridge to professional codes. Ethical systems cannot assume everyone is acting responsibly, but they also cannot treat every mistake as malicious. Good design must handle uncertainty, error, misuse, abuse, incentives, and accountability.

The Professional Problem

How do professionals avoid becoming sources of harm?

Through:



Shift from bad actors in the system to professionals inside the system. The point is not only to identify harmful users or institutions, but also to ask how engineers, moderators, platform designers, and organizations avoid becoming harmful through negligence, opacity, poor design, or reckless deployment.

How do professionals avoid becoming sources of harm?

Professional Codes as Guardrails

Professional codes operationalize ethical priorities as standards of conduct.

They define what responsible practice requires.

They do not eliminate ambiguity.

They structure it.



What This Means

Professional codes of ethics — such as the ACM Code of Ethics and the IEEE Code of Ethics — are formal documents created by professional organizations to translate broad ethical principles into concrete standards of conduct. The metaphor of "guardrails" is deliberate: guardrails do not steer the car, but they keep it on the road. Similarly, professional codes do not resolve every ethical dilemma a practitioner will face, but they establish boundaries, priorities, and expectations that structure decision-making. They operationalize values like fairness, transparency, harm avoidance, and accountability into obligations that professionals can be held to.

The key distinction on this slide is between **eliminating ambiguity** and **structuring it**. Codes do not provide a single correct answer to every situation. Instead, they narrow the range of defensible options, clarify what considerations must be weighed, and create a shared vocabulary for professional responsibility. When a professional faces a gray area — for example, whether to delay a product launch due to unresolved bias in an algorithm — a code does not decide for them, but it tells them what factors must be considered and what cannot be ignored.

Structured Ambiguity — What It Means and Why It Matters

"Structured ambiguity" is the core idea on this slide and deserves careful unpacking. **In ethics, ambiguity is unavoidable** — real-world situations rarely map cleanly onto a single principle or rule. The question is not whether ambiguity exists, but whether professionals face it with or without guidance. Unstructured

ambiguity means every individual improvises: there is no shared framework, no common vocabulary, no agreed-upon priorities. Structured ambiguity means the profession has done the work of identifying what matters, ranking competing values, and establishing processes for reasoning through hard cases — even though the final judgment still belongs to the practitioner.

Think of it this way: a professional code does not say "always choose privacy over transparency." But it does say both privacy and transparency are obligations, it explains why each matters, and it provides a framework for weighing them against each other in context. The ambiguity remains — you still have to decide — but now the decision has structure. You know what values are in play, what trade-offs are legitimate, and what reasoning you need to be able to defend.

This matters because structured ambiguity enables three things that unstructured ambiguity does not: (1) Accountability — if there is a shared framework, a professional can be asked to justify their decision against it, and peers can evaluate whether the reasoning was sound. (2) Consistency — different professionals facing similar dilemmas are working from the same set of priorities, even if they reach different conclusions. (3) Institutional learning — when decisions are made within a structure, organizations can review outcomes, identify patterns, and improve the framework over time.

Concrete example for discussion: A content moderation team at a social media company must decide whether to remove a post that contains medical misinformation but is shared by a licensed physician expressing a minority scientific opinion. An unstructured approach leaves each moderator to their own instincts. A structured approach — guided by a professional code — asks: What are our obligations to public safety? To free expression? To transparency about our decision process? The code does not give the answer, but it ensures the right questions are asked and the reasoning is defensible.

How to Introduce This Concept

This slide is a bridge. The previous slide (The Professional Problem) asked how professionals avoid becoming sources of harm and listed broad strategies — standards, responsibility, disclosure, documentation, testing, transparency, accountability. This slide answers: professional codes are the mechanism that organizes those strategies into a coherent framework.

Suggested introduction: "We just identified what professionals need to do — maintain standards, take responsibility, be transparent. But how do entire professions coordinate around those goals? The answer is professional codes of ethics. These are not abstract philosophical documents. They are operational systems — agreed-upon rules of the road that define what responsible practice looks like in computing and engineering. We are going to look at two of the most important ones: the ACM Code of Ethics and the IEEE Code of Ethics."

Emphasize that codes are living documents maintained by professional communities, not static rules. They evolve as technology and societal expectations change. This sets up the detailed look at ACM and IEEE on the next slides.

ACM Code of Ethics

Association for Computing Machinery · Founded 1947 · World's largest computing society

“Computing professionals’ actions change the world. To act responsibly, they should reflect upon the wider impacts of their work, consistently supporting the public good. The ACM Code of Ethics and Professional Conduct expresses the conscience of the profession.”

KEY PRINCIPLES

Public good

Avoiding harm

Privacy

Fairness

Transparency

Honesty

Competence

Professional responsibility

About ACM

The Association for Computing Machinery (ACM) is the world’s largest educational and scientific computing society, founded in 1947. It serves over 100,000 members across industry, academia, and government worldwide. ACM publishes leading journals and conference proceedings, establishes professional standards, and advocates for computing as a discipline with social responsibilities. Its Code of Ethics and Professional Conduct was first adopted in 1992 and most recently revised in 2018 by an international task force to reflect changes in technology and society.

The Defining Statement

The preamble opens with: “Computing professionals’ actions change the world. To act responsibly, they should reflect upon the wider impacts of their work, consistently supporting the public good. The ACM Code of Ethics and Professional Conduct expresses the conscience of the profession.”

This framing is significant for three reasons. First, it asserts that computing is not a neutral technical activity — it changes the world. This places an affirmative obligation on practitioners. Second, it names the public good as the overriding consideration, not client interest, shareholder value, or institutional loyalty. Third, it describes the code as “the conscience of the profession” — meaning it is not merely a compliance checklist but a moral compass that professionals internalize.

The preamble also states: “The Code is not an algorithm for solving ethical

problems; rather it serves as a basis for ethical decision-making.” This directly connects to the structured ambiguity concept from the previous slide — the code structures judgment without replacing it.

Teaching Notes

Introduce ACM at a high level. Emphasize that this code applies to all computing professionals — not just ACM members — and covers practitioners, instructors, students, and anyone who uses computing technology in an impactful way. Walk through the eight key principles shown on the cards (public good, avoiding harm, privacy, fairness, transparency, honesty, competence, professional responsibility) and note that these are organized into four sections: general ethical principles, professional responsibilities, leadership principles, and compliance. The structure itself reflects the “guardrails” idea — it doesn’t just list values, it organizes them by role and context.

Assignment Framing

This is the point to signal to students that they will be using these codes as analytical tools in their assignment. Ask them to read the quote on the slide carefully. The assignment will require them to apply these principles to a real case — not just identify which principles are relevant, but to reason through how competing principles create structured ambiguity and how the code helps navigate it.

IEEE Code of Ethics

Institute of Electrical and Electronics Engineers · Founded 1963 · 423,000+ members in 160 countries

“We, the members of the IEEE, in recognition of the importance of our technologies in affecting the quality of life throughout the world, and in accepting a personal obligation to our profession, its members and the communities we serve, do hereby commit ourselves to the highest ethical and professional conduct.”

KEY PRINCIPLES

Public safety

Welfare

Sustainability

Technical honesty

Risk disclosure

Competence

Conflicts of interest

Technological
consequences

About IEEE

The Institute of Electrical and Electronics Engineers (IEEE) is the world’s largest technical professional organization, with over 423,000 members in more than 160 countries. It was formed in 1963 through the merger of the American Institute of Electrical Engineers (founded 1884) and the Institute of Radio Engineers (founded 1912). IEEE publishes approximately 200 peer-reviewed journals and over 1,200 conference proceedings annually, and claims to produce over 30% of the world’s literature in electrical, electronics, and computer engineering. Its scope covers engineering and technical systems broadly — not just computing, but hardware, infrastructure, energy, communications, and all technologies that affect quality of life.

The Defining Statement

The IEEE Code opens with: “We, the members of the IEEE, in recognition of the importance of our technologies in affecting the quality of life throughout the world, and in accepting a personal obligation to our profession, its members and the communities we serve, do hereby commit ourselves to the highest ethical and professional conduct.”

This framing differs from ACM’s in important ways. First, it is written as a collective pledge (“We... do hereby commit”) rather than a description of principles — it reads like an oath, not a framework. Second, it explicitly names “quality of life” as the stakes, grounding the code in tangible human outcomes rather than abstract values. Third, it emphasizes “personal obligation” — each engineer accepts individual responsibility, not just organizational compliance.

Where ACM speaks of “the conscience of the profession,” IEEE speaks of a personal commitment to communities served.

Key Contrast with ACM

ACM foregrounds the public good and frames the code as a basis for ethical decision-making. IEEE foregrounds safety, welfare, and sustainability and frames the code as a personal pledge. ACM’s approach is more analytical (“here are principles to reason through”); IEEE’s is more duty-based (“here is what we commit to uphold”). Both are forms of structured ambiguity, but they structure it differently: ACM through deliberative reasoning, IEEE through professional duty and disclosure obligations. Neither is better — they reflect the different traditions of computing (information systems, software) vs. engineering (infrastructure, hardware, safety-critical systems).

Teaching Notes

Introduce IEEE at a high level. Walk through the eight key principles on the cards: public safety, welfare, sustainability, technical honesty, risk disclosure, competence, conflicts of interest, and technological consequences. Note how IEEE’s principles lean toward physical and systemic outcomes — safety, welfare, sustainability, risk — while ACM’s lean toward informational and relational values — privacy, fairness, transparency, honesty. Ask students which framing feels more relevant to AI systems, and why the answer might be “both.”

Assignment Framing

Reinforce that both codes will be tools in the assignment. Students should be able to identify which principles from each code apply to a given case, explain where the two codes converge and diverge, and articulate how the structured ambiguity in each code shapes their analysis differently. The assignment is not about picking the “right” code but about using both as complementary lenses.

ACM / IEEE Divergence Example

Same system. Different ethical emphasis.

**ACM may foreground
Fairness · Non-
discrimination ·
Transparency · Avoidable
harm**

vs.

**IEEE may foreground
Safety · Technical
reliability · Risk disclosure
· Insufficient validation**

Framing Notes: Use this as the concrete example of possible divergence. Both frameworks may criticize the deployment, but not with the same center of gravity. ACM can foreground inequity and avoidable harm to disadvantaged users. IEEE can foreground engineering risk, insufficient testing, and reckless deployment of an unreliable system.

Full Slide Text (Addendum):
Same system. Different ethical emphasis.

Scenario: A COVID misinformation classifier has a higher false-positive rate for certain dialect communities. The company deploys it anyway to reduce harmful false claims quickly.

ACM may foreground: Fairness, Non-discrimination, Transparency, Avoidable harm

IEEE may foreground: Safety, Technical reliability, Risk disclosure, Insufficient validation

EU Governance Framing

EU AI Act (Regulation 2024/1689) · Adopted June 2024 · Institutional regulatory framework

“The purpose of this Regulation is to... promote the uptake of human-centric and trustworthy artificial intelligence, while ensuring a high level of protection of health, safety, fundamental rights... including democracy, the rule of law and environmental protection, against the harmful effects of AI systems.” — EU AI Act, Article 1

KEY PRINCIPLES

Oversight

Proportionality

Transparency

Accountability

Rights

Systemic risk

Public-interest governance

About the EU AI Act

The EU AI Act (Regulation 2024/1689) was adopted by the European Parliament and Council on 13 June 2024, making it the world’s first comprehensive legal framework for artificial intelligence. It entered into force on 1 August 2024, with provisions phasing in through 2027. It builds on earlier work by the EU High-Level Expert Group on AI, which in 2019 published Ethics Guidelines for Trustworthy AI establishing that trustworthy AI must be (1) lawful, (2) ethical, and (3) robust. The AI Act translates these aspirations into binding law, with enforcement mechanisms, risk classifications, and institutional oversight structures.

The Defining Statement

Article 1 states: “The purpose of this Regulation is to improve the functioning of the internal market and promote the uptake of human-centric and trustworthy artificial intelligence (AI), while ensuring a high level of protection of health, safety, fundamental rights enshrined in the Charter, including democracy, the rule of law and environmental protection, against the harmful effects of AI systems in the Union and supporting innovation.”

Notice the fundamentally different framing from the professional codes. ACM and IEEE speak to practitioners: what should you, as a professional, do? The EU AI Act speaks to systems and institutions: what must be ensured before an AI system can be placed on the market? The subject is not the engineer’s conscience but the regulatory apparatus — market surveillance, conformity assessment, risk classification, and enforcement.

Governance Models vs. Professional Codes — The Core Contrast

This is the critical distinction to draw on this slide. Professional codes and governance frameworks both address ethical responsibility, but they operate at different levels and through different mechanisms:

1. Authority: Professional codes are adopted by voluntary membership organizations. Governance frameworks are enacted by legislatures and enforced by regulatory bodies. A code appeals to professional conscience; a regulation compels compliance under penalty of law.
2. Subject: Codes address individual practitioners (“a computing professional should...”). Governance frameworks address systems, organizations, and markets (“AI systems placed on the market must...”).
3. Mechanism: Codes structure ambiguity through principles and values — they guide reasoning. Governance frameworks structure ambiguity through risk categories, compliance requirements, and institutional oversight — they constrain action.
4. Enforcement: Violation of a professional code may result in loss of membership or peer censure. Violation of a governance framework can result in fines, market bans, and legal liability. The EU AI Act can impose fines up to €35 million or 7% of global annual turnover.
5. Scope: Codes apply to professionals within a discipline. Governance frameworks apply to anyone who develops, deploys, or distributes AI systems in a jurisdiction — regardless of professional affiliation.

Teaching Notes

Frame this slide as a shift in level. The previous two slides asked: what should individual professionals do? This slide asks: what should institutions, markets, and governments ensure? Walk through the seven principles on the cards — oversight, proportionality, transparency, accountability, rights, systemic risk, public-interest governance — and note how they differ from the ACM/IEEE principles. These are not virtues a practitioner cultivates; they are structural requirements an institution must satisfy. The EU AI Act does not ask whether an engineer acted in good conscience. It asks whether a system was properly classified, assessed, monitored, and documented before deployment.

Assignment Framing

Students now have three lenses for their assignment: ACM (professional conscience, public good), IEEE (personal obligation, safety and welfare), and the EU governance framework (institutional oversight, rights protection, systemic risk). The assignment should require them to show how the same case looks different through each lens — not just which principles apply, but what kind of reasoning each framework demands and what it leaves unresolved. The governance lens adds a layer the codes cannot: who is responsible when the system itself is the problem, not the individual who built it?

Three Levels of Ethical Structure

Professional ethics

What should responsible practitioners do?

Governance ethics

How should institutions manage risk and legitimacy?

Mediation / power analysis

Who can observe, control, and authorize action?

Make the distinction explicit: ACM and IEEE focus on professional conduct; EU-oriented materials focus on governance and public accountability; BBS/EDOCA analyzes observability, control, authority, and responsibility distribution.

Ethical Frameworks Identify Values

Frameworks identify values.

Values ask:

What matters?

BBS / EDOCA identifies power structure.

Power asks:

Who can act?

Who can see?

Who can authorize?

Who is affected without recourse?

This slide draws a key distinction between what BBS/EDOCA does and what ethical frameworks do. They are complementary, not redundant.

BBS/EDOCA is an analytical tool. It helps you map out the structural landscape of a system: Where is the risk? Where can harm occur? Who holds authority? Who has access? Who can observe what is happening? Who is affected and has no recourse? These are descriptive and structural questions. BBS does not tell you whether something is right or wrong — it tells you where to look.

Ethical frameworks — such as ACM, IEEE, EU governance principles, or classical normative theories — are the ones that apply a specific set of values to those areas. They say: given that this person holds authority, here is what responsible use of that authority looks like. Given that harm can occur here, here is what counts as acceptable risk. Frameworks supply the evaluative lens; BBS supplies the map.

The practical takeaway: you need both. BBS without a framework identifies structure but cannot judge it. A framework without BBS applies values but may miss where power, access, and harm actually sit in the system. Together, they produce a grounded ethical analysis.

THIS IS WHERE WE STOPPED

I PLAYED THE COVID VIDEO AND IT WAS 10 minutes over

GROUP TASK

Return to the expanded COVID case.

Identify:

Stakeholders

Benefits

Harms

Kinds of actors

Observability

Control

Authority

Framework conflicts

Unresolved ambiguity

All groups use the same expanded COVID case study. The shared case is fixed so framework comparisons across groups remain meaningful.

Actor Analysis

In the case, identify examples of:

Misinformed actors

Negligent actors

Reckless actors

Strategic actors

Malicious actors

Institutional actors

Procedural actors

Students are still working from the expanded COVID case, not the short opening card.

Framework Comparison

How would each framework evaluate the same COVID case?

ACM	IEEE	EU governance	BBS / EDOCA
<i>Where do they agree?</i>	<i>Where do they diverge?</i>	<i>What does each make visible?</i>	<i>What does each under-describe?</i>

Keep the case fixed. The comparison is meaningful because every group is evaluating the same moderation system.

Observability, Control, Authority

Who can see the decision?

Who can change the system?

Who can authorize action?

Who is affected but cannot intervene?

Use the expanded COVID case here. Public-health authorities may have expertise, platforms may have operational control, regulators may have formal authority, and users may have limited observability and limited recourse.

Ethical Ambiguity

The question is not:

Was moderation simply right or wrong?

The question is:

How do different frameworks assign responsibility, define harm, and evaluate the same system differently?

Prevent premature closure. Students are not being asked to arrive at a final moral verdict. They are being asked to demonstrate structured ethical analysis.

Closing Synthesis

Ethical systems must account for:

-  Good-faith error
-  Bad-faith manipulation
-  Professional responsibility
-  Institutional incentives
-  Governance limits
-  Framework disagreement
-  Mediation pathways
-  Power asymmetry

Name the main takeaway. Modern socio-technical systems cannot rely only on personal virtue or good intentions. They require professional standards, governance structures, and mediation-aware analysis.

Final Questions

How should systems respond when they can observe a false claim but not the intent behind it?

How do professional codes help prevent practitioners from becoming sources of harm?

Why might ACM and IEEE reach different evaluations of the same system?

Use these as closing discussion prompts. If time is short, choose one. The first question is directly connected to the COVID case. The second transitions to ACM/IEEE reading. The third prepares the follow-up assignment.

TRANSITION TO ASSIGNMENT: If we get here, use this slide as the pivot. Say something like: “These questions are exactly what the assignment asks you to work through. We may not have time to walk through the full assignment now, but the remaining slides frame it. You can look at them on your own and see how they set up what you need to do.”

If time permits, move to the next slide for the assignment preview. If not, tell students the assignment is posted and these final slides give them the framing they need.

Follow-Up Assignment Preview

ACM and IEEE Professional Ethics Applied to the COVID Case

Read	Apply them to	Compare
• ACM Code of Ethics • IEEE Code of Ethics	• One intervention strategy from the shared COVID misinformation case	• Ethical emphasis • Relevant principles • Professional responsibility • Harms identified • What remains unclear

ASSIGNMENT OVERVIEW: This is a one-week assignment, due Week 12, 800–1200 words. The scope has been narrowed to focus on professional ethics only — ACM and IEEE applied to one intervention strategy from the shared COVID misinformation case. No BBS/EDOCA analysis, no EU governance comparison, no full stakeholder mapping.

WHAT STUDENTS DO: Choose one intervention strategy from the expanded case (labeling, downranking, removing, redirecting, suspending repeat violators, automated moderation, or recommendation system design). Briefly identify the governance context. Then apply two ACM principles and two IEEE principles to that strategy. Compare how ACM and IEEE differ in emphasis, scope, and what they make visible. Close with a short reflection on which code gives more useful guidance for that specific case.

THE FIVE PARTS: (1) Intervention Strategy and Case Focus — what strategy, who designs/maintains it, who is affected, why it is ethically complicated. (2) ACM Code Analysis — two principles applied to the case. (3) IEEE Code Analysis — two principles applied to the case. (4) Comparison — where they agree, where they diverge, how differences affect evaluation. A comparison table is suggested. (5) Reflection — which code gives more useful guidance for this specific case and why.

KEY FRAMING: The central question is “How do ACM and IEEE define responsible professional conduct differently when applied to the same misinformation intervention strategy?” Students are not deciding whether COVID moderation

was right or wrong. They are comparing how two professional codes apply their principles. Keep attention on professional ethical responsibility, not general political debate.

STRONG VS WEAK WORK: Strong work picks a specific intervention, selects relevant principles (not just summarizes the codes), and applies them directly to professional conduct in the case. Weak work discusses COVID misinformation in general without identifying a specific professional action, or treats ACM and IEEE as interchangeable.

HOW THE SLIDES FRAME THIS: If you did not get to this slide in class, tell students to review slides 37–42 on their own. The framework comparison slide (37) shows the broader landscape of ACM, IEEE, EU governance, and BBS/EDOCA. The observability/control/authority slide (38) gives them structural vocabulary. The ethical ambiguity slide (39) sets the right question. The final questions slide (41) gives them the exact kind of thinking the assignment requires. This preview slide (42) maps the assignment structure directly.

EXIT REFLECTION

In your private journal:

What changed in your ethical evaluation after introducing bad actors and professional codes?

What did your first answer fail to consider?

Return students to individual reflection. Reinforce metacognition and create continuity with prior class practice. The instructor may ask for voluntary sharing if time remains, but the journal response can remain private.