

SEID 2364 – Assignment 4 Ethical Frameworks and the Evaluation of Harm

Part 1 – System Selection and Harm Recap

Selected System: AI Hiring System

An AI hiring system is used by companies to evaluate job applicants and recommend candidates for employment. The system processes resumes, qualifications, experience, and other applicant data to assist recruiters in making hiring decisions. Companies use these systems to reduce hiring time and improve efficiency.

However, the system can create several harms. If historical hiring data contains bias, the AI may unfairly favor certain groups while disadvantaging others. This may result in discrimination based on gender, ethnicity, or educational background. Applicants may also experience economic harm if they are unfairly denied opportunities. Another concern is transparency because applicants often cannot fully understand how the system reached its decision.

Part 2 – Core Ethical Framework

Selected Framework: Consequentialism (Utilitarianism)

Consequentialism is an ethical framework that evaluates actions based on their outcomes or consequences. The most common form is utilitarianism, which focuses on maximizing overall happiness or benefit while reducing harm. Philosophers such as Jeremy Bentham and John Stuart Mill are strongly connected to this framework.

In the context of AI hiring systems, consequentialism would evaluate whether the system produces more benefits than harms. If the system improves efficiency, reduces costs, and helps companies process applications faster, the framework may view it positively. However, if the system creates unfair discrimination or harms applicants, those negative consequences must also be considered.

Consequentialism prioritizes the overall outcome rather than individual intentions. This means that even if a company intends to improve hiring, the system may still be unethical if the consequences produce significant social harm.

Part 3 – Independent Framework Research

Selected Framework: Deontology

Deontology is an ethical framework associated with philosopher Immanuel Kant. It focuses on duties, rules, and moral obligations rather than consequences. According to deontology, actions are ethical when they follow moral principles and respect individuals as ends in themselves rather than as tools.

In AI hiring systems, deontology would emphasize fairness, equal treatment, and respect for applicants. A hiring system that discriminates against certain groups would be unethical even if it increases efficiency because it violates moral duties of fairness and equality.

Framework Evaluation

Deontology clearly defines harm in terms of violating rights or duties. It is useful because it protects individuals from being treated unfairly. However, the framework can become difficult when rules conflict. For example, a company may want efficiency while also ensuring fairness. Deontology may also struggle in situations where strict rules reduce overall benefits.

One questionable scenario is when a company uses an AI system that improves hiring speed for thousands of applicants but unintentionally disadvantages a small group. Consequentialism may accept the trade-off if overall benefits are higher, while deontology would reject the unfair treatment regardless of efficiency.

Part 4 – Multi-Framework Harm Evaluation

Consequentialism would focus on whether the AI hiring system creates greater benefits than harms. It may support the system if it improves efficiency, reduces hiring costs, and increases productivity. However, it would criticize the system if discriminatory harms become widespread.

Deontology would focus on fairness and equal treatment. It would prioritize protecting applicants from discrimination, even if doing so reduces efficiency or increases hiring costs. This framework values moral duties above overall outcomes.

Part 5 – Comparative Analysis

Both frameworks recognize that discrimination in hiring is harmful. However, they differ in how they evaluate trade-offs.

Consequentialism focuses on overall outcomes and may accept some harm if larger benefits exist. Deontology focuses on duties and rights, meaning unfair treatment is unacceptable even if the overall system benefits.

Deontology reveals harms connected to fairness and rights more strongly, while consequentialism highlights social and economic outcomes. Both frameworks are useful, but each leaves gaps. Consequentialism may justify unfair outcomes, while deontology may ignore practical consequences.

Part 6 – Meta-Evaluation

No single ethical framework is fully sufficient for analyzing complex AI systems. AI hiring systems involve multiple stakeholders, competing interests, and different forms of harm. Consequentialism helps evaluate outcomes, while deontology protects fairness and rights.

I believe combining multiple ethical frameworks provides a stronger ethical analysis because it allows both consequences and moral duties to be considered together. This creates a more balanced approach to evaluating harm within AI systems.